

Importing Libraries

```
In [528...  
import os  
os.getcwd()
```

```
Out[528... 'C:\\BigData\\~notebookJupyter'
```

```
In [529...  
import requests  
requests.packages.urllib3.disable_warnings()
```

```
In [530...  
import plotly.plotly as py  
import plotly.graph_objs as go
```

```
In [531...  
import findspark  
import pandas as pd  
findspark.init("C:\\BigData\\Spark")  
findspark.find()
```

```
Out[531... 'C:\\BigData\\Spark'
```

```
In [532...  
import pyspark  
from pyspark import SparkContext  
from pyspark import SQLContext  
from pyspark.sql import SparkSession  
from pyspark.sql.functions import *
```

```
In [533...  
spark.conf.set("spark.sql.execution.arrow.enabled", "true")
```

```
In [534...  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [535...  
sc = SparkContext.getOrCreate()  
spark=SparkSession.builder.getOrCreate()  
print(sc.version)
```

2.4.7

```
In [536...  
from pyspark.sql.types import IntegerType
```

```
In [537...  
plt.rcParams.update({'figure.figsize':(10,8), 'figure.dpi':100})
```

Loading Dataset

```
In [538...  
%%time  
orig_df = pd.read_csv("C:/Users/HP/Desktop/ECE552_final/ECE552_project/BitcoinHeistD  
  
Wall time: 3.06 s
```

```
In [607... orig_df.head()
orig_df.reset_index(drop=True, inplace=True)
```

converting into bitcoin

```
In [540... orig_df['income'] = orig_df['income']/100000000
```

```
In [541... orig_df=orig_df.dropna()
```

```
In [542... orig_df.to_csv("C:/Users/HP/Desktop/ECE552_final/ECE552_project/BitcoinHeistData1.csv")
```

reading data using spark

```
In [543... %%time
df = spark.read.format("csv"). \
    option("header", "true").option("mode", "DROPMALFORMED").option("delimiter", ",") \
    option("ignoreLeadingWhiteSpace", "true").option("ignoreTrailingWhiteSpace", "true") \
    option("inferSchema", "true"). \
    load('C:/Users/HP/Desktop/ECE552_final/ECE552_project/BitcoinHeistData1.csv')
```

Wall time: 3.98 s

```
In [544... print(df.count())
df.printSchema()
```

```
1048575
root
 |-- _c0: integer (nullable = true)
 |-- address: string (nullable = true)
 |-- year: integer (nullable = true)
 |-- day: integer (nullable = true)
 |-- length: integer (nullable = true)
 |-- weight: double (nullable = true)
 |-- count: integer (nullable = true)
 |-- looped: integer (nullable = true)
 |-- neighbors: integer (nullable = true)
 |-- income: double (nullable = true)
 |-- label: string (nullable = true)
```

Writing parquet

```
In [545... df.repartition(4).write.mode('overwrite').parquet('C:/Users/HP/Desktop/ECE552_final/BitcoinHeistData1.parquet')
```

```
In [546... %%time
spark=SparkSession.builder.appName("Python Spark dataframe").config("spark.some.config", "true").getOrCreate()
parquet_df=spark.read.parquet('C:/Users/HP/Desktop/ECE552_final/ECE552_project/BitcoinHeistData1.parquet')
```

Wall time: 146 ms

```
In [547... type(parquet_df)
```

```
Out[547... pyspark.sql.dataframe.DataFrame
```

In [548...

parquet_df.show()

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
|  _c0|          address|year|day|length|          weight|count|looped|neighbors|
income|          label|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
|207249|1EkNTexkRsMpgTE7a...|2011|176|    124|    2.98E-8|    2|    0|    2|
2.25|          white|
|121548|13JrBxwSJ EjXtvAwa...|2011| 90|     6|    0.75|    3|    0|    2|
80.0|          white|
| 97738|18CRxe48uz1a3LWgd...|2011| 67|     0|    0.5|    1|    0|    2|
116.91|          white|
| 75564|17hFuRueNZAiPwDr1...|2011| 44|    16|  0.03125|    1|    0|    1|
0.3|          white|
| 72434|14XRZVVCi6j7aH6U9...|2011| 41|    22| 2.44141E-4|    1|    0|    2|
0.5|          white|
|210188|1JucwoRuVqAh4g7Lu...|2011|179|     0|    0.5|    1|    0|    1|
2.93118|          white|
| 84220|1Q4WQ7irizHbkt2fM...|2011| 53|    58|0.506404332|   51|    0|    2|
48.76|          white|
|132915|19L7TurgXprUJpN4v...|2011|102|    30|  1.2207E-4|    1|    0|    1|
58.38|          white|
|195332|1Pg3JkPdmTeSjLxJk...|2011|164|   114|   8.38E-9|    2|    0|    3|
2.92|          white|
| 93103|1AYZQFjKDvETaUP8w...|2011| 62|    16|0.001953125|    1|    0|    2|
11.33|          white|
|266715|153SRUhhVpfmnat3h...|2011|236|     6|    0.5|    1|    0|    2|
0.6|          white|
|137696|1271tY1rX2ihMocWF...|2011|106|     0|    0.5|    1|    0|    1|
1.01|          white|
|217213|1NDGzwDu9ADJdwoC1...|2011|186|    60|   1.53E-5|    1|    0|    1|
0.43|          white|
|161057|1ApE5DL1pL5NTqfBb...|2011|130|     0|    1.0|    1|    0|    2|
6.88|          white|
| 50877|1H8EvrtmCK6GUh8MN...|2011| 16|     6|    0.25|    1|    0|    2|
45.49|          white|
|275645|187rbGVNVw6MR5GPn...|2011|244|     2|    0.5|    1|    0|    2| 47.
50271007|          white|
|232600|1D88h6L5d47eqCwGX...|2011|201|    50|   4.11E-6|   28|    0|    1|
1.0|          white|
|162132|1BK47BchauVHBHCj...|2011|131|     0|    1.0|    1|    0|    2|
2.11|          white|
|206442|1417elt5Xu9GhqHxN...|2011|175|    34| 0.00390625|    1|    0|    2| 563.
01393392|          white|
| 19473|1CweHaxHBhk7pt6VZ...|2015|265|   144| 0.11639657| 3544| 1115|    1|
3.13|paduaCryptoWall|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
only showing top 20 rows

```

In [549...

```

print(parquet_df.count())
parquet_df.printSchema()

```

```

1048575
root
|-- _c0: integer (nullable = true)
|-- address: string (nullable = true)
|-- year: integer (nullable = true)
|-- day: integer (nullable = true)

```

```
-- length: integer (nullable = true)
-- weight: double (nullable = true)
-- count: integer (nullable = true)
-- looped: integer (nullable = true)
-- neighbors: integer (nullable = true)
-- income: double (nullable = true)
-- label: string (nullable = true)
```

In [550...

```
#import org.apache.spark.sql.types.IntegerType
#bitcoin_df["income"].astype("int64")
```

In [551...

```
print(parquet_df.schema["income"].dataType)
```

DoubleType

In [552...

```
parquet_df.select([count(when(col(c).isNull(), c)).alias(c) for c in parquet_df.colu
```

```
+---+-----+-----+---+-----+-----+-----+-----+-----+-----+-----+
|_c0|address|year|day|length|weight|count|looped|neighbors|income|label|
+---+-----+-----+---+-----+-----+-----+-----+-----+-----+-----+
| 0|      0|  0| 0|    0|    0|  0|    0|        0|    0|    0|
+---+-----+-----+---+-----+-----+-----+-----+-----+-----+-----+

```

In [553...

```
parquet_df = parquet_df.na.drop()
```

In [554...

```
parquet_df.show()
```

```
+-----+-----+-----+---+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+
|  _c0|          address|year|day|length|    weight|count|looped|neighbors|
income|          label|
+-----+-----+-----+---+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+
|207249|1EkNTErkRsMpgTE7a...|2011|176|  124|  2.98E-8|  2|  0|    2|
2.25|      white|
|121548|13JrBxWSJEjXtvAwa...|2011| 90|   6|    0.75|  3|  0|    2|
80.0|      white|
| 97738|18CRxe48uz1a3LWgd...|2011| 67|   0|    0.5|  1|  0|    2|
116.91|      white|
| 75564|17hFuRueNZAiPwDr1...|2011| 44|  16|  0.03125|  1|  0|    1|
0.3|      white|
| 72434|14XRZVVCi6j7aH6U9...|2011| 41|  22| 2.44141E-4|  1|  0|    2|
0.5|      white|
|210188|1JucwoRuVqAh4g7Lu...|2011|179|   0|    0.5|  1|  0|    1|
2.93118|      white|
| 84220|1Q4WQ7irizHbkt2fM...|2011| 53|  58|0.506404332| 51|  0|    2|
48.76|      white|
|132915|19L7TurgXprUjPn4v...|2011|102|  30| 1.2207E-4|  1|  0|    1|
58.38|      white|
|195332|1Pg3JkPdmTeSjLxJk...|2011|164| 114|  8.38E-9|  2|  0|    3|
2.92|      white|
| 93103|1AYZQFjKDvETaUP8w...|2011| 62|  16|0.001953125|  1|  0|    2|
11.33|      white|
|266715|153SRUhhVpfmnat3h...|2011|236|   6|    0.5|  1|  0|    2|
0.6|      white|
|137696|1271tY1rX2ihMocWF...|2011|106|   0|    0.5|  1|  0|    1|
1.01|      white|

```

```

|217213|1NDGzwDu9ADJdwoC1...|2011|186|    60|    1.53E-5|    1|    0|    1|
0.43|    white|
|161057|1ApE5DL1pL5NTqfBb...|2011|130|    0|    1.0|    1|    0|    2|
6.88|    white|
| 50877|1H8EvrtmCK6GUh8MN...|2011| 16|    6|    0.25|    1|    0|    2|
45.49|    white|
|275645|187rbGVNVw6MR5GPn...|2011|244|    2|    0.5|    1|    0|    2| 47.
50271007|    white|
|232600|1D88h6L5d47eqCwGX...|2011|201|   50|   4.11E-6|   28|    0|    1|
1.0|    white|
|162132|1BK47BchauWVHBHCj...|2011|131|    0|    1.0|    1|    0|    2|
2.11|    white|
|206442|1417eLt5Xu9GhqHxN...|2011|175|   34| 0.00390625|    1|    0|    2|563.
01393392|    white|
| 19473|1CweHaxHBhk7pt6VZ...|2015|265|  144| 0.11639657| 3544|  1115|    1|
3.13|paduaCryptoWall|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+

```

only showing top 20 rows

In []:

In [555... `from pyspark.sql.functions import desc`

In [556... `parquet_df.createOrReplaceTempView("bc_data")`

In [557... `income_year = spark.sql("SELECT income,year from bc_data order by income desc")`

Highest transaction from the data set and the year it occurred

In [558... `income_year.head(3) ##499643 will be highest`

Out[558... `[Row(income=499643.98238996, year=2011),
Row(income=498244.66601007, year=2011),
Row(income=497976.31116927, year=2011)]`

The Year wise transation count

In [559... `import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns`

In []:

In [560... `Year_txn=spark.sql("select year,count(year) as number_of_txn_per_year from bc_data w`

In [561... `Year_txn.show()`

```

+-----+-----+
|year|number_of_txn_per_year|
+-----+-----+

```

2012	365714
2011	355349
2013	294372
2014	10319

+-----+-----+

In [562... `type(Year_txn)`

Out[562... `pyspark.sql.dataframe.DataFrame`

In [563... `Year_txn=Year_txn.toPandas()`

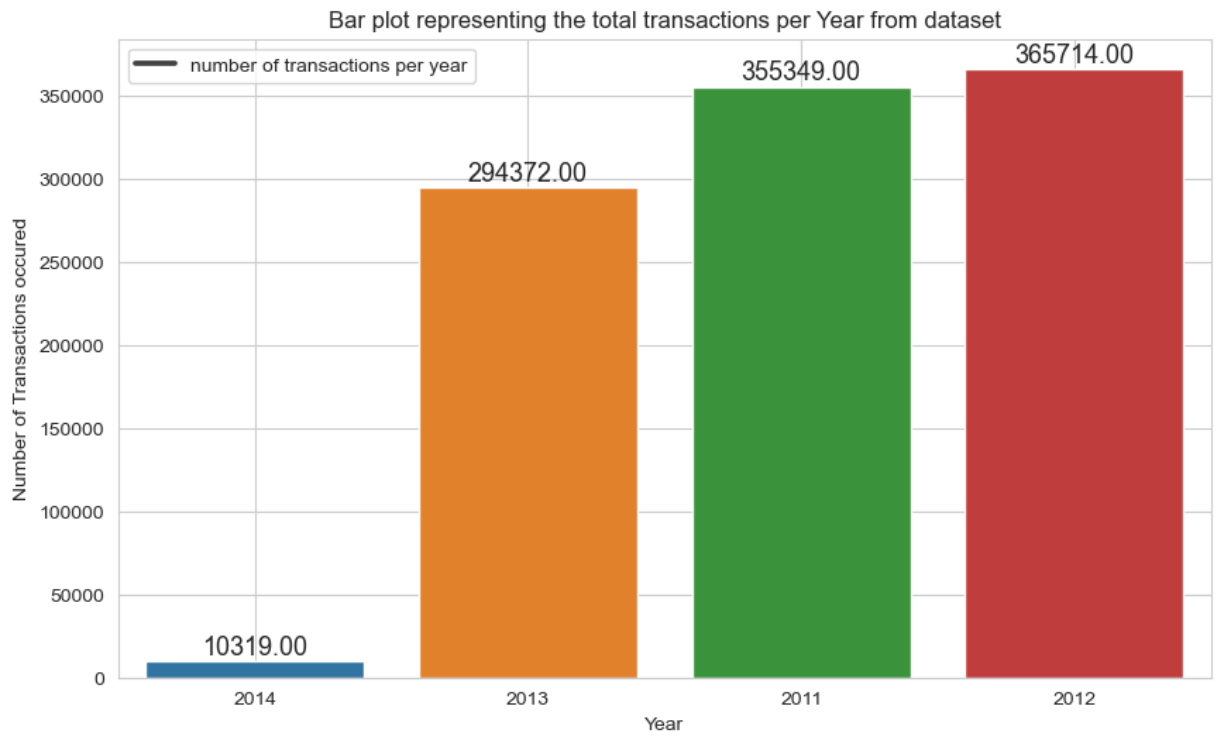
In [564... `temp1=Year_txn["year"]`
`temp=Year_txn["number_of_txn_per_year"]`

In [565... `plt.figure(figsize=(10, 6))`
`plots = sns.barplot(x="year", y="number_of_txn_per_year", data=Year_txn, order=Year_t`

Iterating over the bars one-by-one
`for bar in plots.patches:`

`plots.annotate(format(bar.get_height(), '.2f'),`
`(bar.get_x() + bar.get_width() / 2,`
`bar.get_height()), ha='center', va='center',`
`size=13, xytext=(0, 7),`
`textcoords='offset points')`

Creating the legend of the bars in the plot
`plt.legend(labels = ['number of transactions per year'])`
Giving the tilte for the plot
`plt.title("Bar plot representing the total transactions per Year from dataset")`
Naming the x and y axis
`plt.xlabel('Year')`
`plt.ylabel('Number of Transactions occured')`
`plt.grid(True)`



In []:

The Count of labels

In [566..

```
#minned minned label
parquet_df.select("label").distinct().show(truncate=False)
```

```
+-----+
|label|
+-----+
|paduaJigsaw|
|montrealXLockerv5.0|
|montrealJigSaw|
|montrealCryptoTorLocker2015|
|montrealXTPLocker|
|montrealAPT|
|montrealRazy|
|white|
|montrealNoobCrypt|
|princetonCerber|
|montrealEDA2|
|paduaCryptoWall|
|montrealGlobeImposter|
|montrealCryptConsole|
|montrealSamSam|
|montrealSam|
|montrealDMALockerv3|
|montrealXLocker|
|montrealVenusLocker|
|montrealCryptoLocker|
+-----+
only showing top 20 rows
```

In [567..

```
label_df = spark.sql("select label, count(label) as label_count from bc_data group b
```

```
In [568... label_df=label_df.na.drop()
```

```
In [569... Newlabel_df=label_df.toPandas()
```

```
In [570... Newlabel_df=Newlabel_df.head(6)
Newlabel_df
```

```
Out[570...
      label  label_count
0      white      1007162
1  paduaCryptoWall      12390
2  montrealCryptoLocker      9315
3   princetonCerber      9223
4   princetonLocky      6625
5  montrealCryptXXX      2419
```

White(1007162) which is not known to be ransomware

```
In [571... print(type(label_df))
print(type(Newlabel_df))
x=Newlabel_df['label']
y=Newlabel_df["label_count"]

<class 'pyspark.sql.dataframe.DataFrame'>
<class 'pandas.core.frame.DataFrame'>
```

```
In [572... Newlabel_df=Newlabel_df.drop([0])
x=Newlabel_df['label']
y=Newlabel_df["label_count"]
Newlabel_df
```

```
Out[572...
      label  label_count
1  paduaCryptoWall      12390
2  montrealCryptoLocker      9315
3   princetonCerber      9223
4   princetonLocky      6625
5  montrealCryptXXX      2419
```

```
In [573... import numpy as np

l = Newlabel_df['label']
data = Newlabel_df["label_count"]

# Creating explode data
explode = (0.1, 0.0, 0.2, 0.3, 0.0)

# Creating color parameters
colors = ( "orange", "cyan", "brown",
```



```

        "grey", "indigo", "beige")

# Wedge properties
wp = { 'linewidth' : 1, 'edgecolor' : "green" }


# Creating autocpt arguments
def func(pct, allvalues):
    absolute = int(pct / 100.*np.sum(allvalues))
    #print(absolute)
    #return pct
    return "{:.1f}%\n({:d})".format(pct, absolute)

# Creating plot
fig, ax = plt.subplots(figsize =(10, 8))
wedges, texts, autotexts = ax.pie(y,
                                   autopct = lambda pct: func(pct, data),
                                   explode = explode,
                                   labels = x,
                                   shadow = True,
                                   colors = colors,
                                   startangle = 105,
                                   wedgeprops = wp,
                                   textprops = dict(color ="red"))

# Adding Legend
ax.legend(wedges, l,
          title ="Labels",
          loc ="center left",
          bbox_to_anchor =(1, 0, 0.5, 1))

plt.setp(autotexts, size = 11.5, weight ="bold")
ax.set_title("Labels known to be ransomware")
# show plot
plt.show()

```



```
In [574... #plt.pie(NewLabel_df['label'],labels = NewLabel_df['label_count'])
```

```
In [575... income_year = spark.sql("SELECT income , year from bc_data where year = 2011 or year
```

```
In [576... income_year.show()
```

```
+-----+-----+
|      income|year|
+-----+-----+
|499643.98238996|2011|
|498244.66601007|2011|
|497976.31116927|2011|
|497852.01883234|2011|
| 496785.7983011|2011|
|496766.69306791|2011|
|496651.52076174|2011|
|496324.31999242|2011|
|496150.15001557|2011|
|486752.21419529|2011|
|466023.66865338|2011|
|464548.09728703|2011|
|455501.30057744|2011|
|454856.92132679|2011|
|451523.24401692|2011|
|449111.17551607|2011|
|447785.59485753|2011|
|447427.86485753|2011|
|446573.52271277|2011|
|442561.65706485|2011|
+-----+-----+
only showing top 20 rows
```

```
In [ ]:
```

```
In [577... income_year=income_year.toPandas()
```

```
In [578... income_year.dtypes
```

```
Out[578... income    float64
year         int32
dtype: object
```

```
In [579... income_year.head(5)
```

```
Out[579...      income  year
0  499643.982390  2011
1  498244.666010  2011
2  497976.311169  2011
3  497852.018832  2011
4  496785.798301  2011
```

In [580...

income_year.astype(int)

Out[580...

	income	year
0	499643	2011
1	498244	2011
2	497976	2011
3	497852	2011
4	496785	2011
...
649716	0	2013
649717	0	2013
649718	0	2013
649719	0	2013
649720	0	2013

649721 rows × 2 columns


In [581...

income_year=income_year[income_year.income != 0]

boxplot

In [582...

ax = sns.boxplot(x="year", y="income", data=income_year)



1/7/22, 1:38 PM

ece552project_team[14]

In []:

ax = sns.swarmplot(x="year", y="income", data=income_year, color=".25")

In []:

In [583...]

orig_df1 = orig_df[orig_df.length != 0]

In [584...]

orig_df1.reset_index(drop=True, inplace=True)

In [585...]

orig_df1.to_csv("C:/Users/HP/Desktop/ECE552_final/ECE552_project/BitcoinHeistData2.c

In [586...]

house_df = spark.read.format("csv"). \
option("header", "true").option("mode", "DROPMALFORMED").option("delimit
option("ignoreLeadingWhiteSpace", "true").option("ignoreTrailingWhiteSpac
option("inferschema", "true"). \
load('C:/Users/HP/Desktop/ECE552_final/ECE552_project/BitcoinHeistData1.

In []:

In []:

In [587...]

house_df.take(1)

Out[587...]

[Row(_c0=0, address='111K8kZAEJg245r2cM6y9zgJGHZtJPY6', year=2017, day=11, length=1
8, weight=0.008333333, count=1, looped=0, neighbors=2, income=1.0005, label='princet
onCerber')]

descriptive analysis

In [588...]

house_df.describe().toPandas().transpose()

Out[588...]

	0	1	2	3
summary	count	mean	stddev	min
_c0	1048575	524287.0	302697.6735952888	0
address	1048575	None	None	111111111111111111111111111111114oLvT2 3GbjoRZ
year	1048575	2012.0483856662613	1.005056531238343	2011
day	1048575	173.52949383687385	100.14085576953202	1
length	1048575	26.317415540137805	42.720042828726775	0
weight	1048575	0.5406222223650033	2.9908459496285897	3.61E-94
count	1048575	158.20701237393607	602.4127303067922	1
looped	1048575	36.08665283837589	288.47339538256995	0
neighbors	1048575	2.3242200128746155	18.623586560640728	1
income	1048575	90.26239897843372	2555.8505269303614	0.3

	0	1	2	3
label	1048575	None	None	montrealAPT

In [589...


```
import pandas as pd
from pandas.plotting import scatter_matrix
```

In [590...

```
numeric_features = [t[0] for t in house_df.dtypes if t[1] == 'int' or t[1] == 'doubl
sampled_data = house_df.select(numeric_features).sample(False, 0.8).toPandas()
```

In [608...

```
axs = scatter_matrix(sampled_data, figsize=(10, 10))
#ax.set_title("Scatter matrix for linear correlation between multiple independent va
n = len(sampled_data.columns)
for i in range(n):
    v = axs[i, 0]
    v.yaxis.label.set_rotation(0)
    v.yaxis.label.set_ha('right')
    v.set_yticks(())
    h = axs[n-1, i]
    h.xaxis.label.set_rotation(90)
    h.set_xticks(())
```



Machine learning models

Linear regression model

```
In [592... from pyspark.ml.feature import VectorAssembler
```

```
In [593... house_df.printSchema()
```

```
root
|-- _c0: integer (nullable = true)
|-- address: string (nullable = true)
|-- year: integer (nullable = true)
|-- day: integer (nullable = true)
|-- length: integer (nullable = true)
|-- weight: double (nullable = true)
|-- count: integer (nullable = true)
|-- looped: integer (nullable = true)
|-- neighbors: integer (nullable = true)
|-- income: double (nullable = true)
|-- label: string (nullable = true)
```

```
In [609...
```

```
import six
for i in house_df.columns:
    if not( isinstance(house_df.select(i).take(1)[0][0], six.string_types)):
        print( "Correlation to Length for ", i, house_df.stat.corr('length',i))
```

```
Correlation to Length for _c0 0.046628208681509936
Correlation to Length for year 0.09177169340486724
Correlation to Length for day 0.04342247835202402
Correlation to Length for length 1.0
Correlation to Length for weight 0.060180331529211345
Correlation to Length for count 0.6627173924217697
Correlation to Length for looped 0.32688044270582417
Correlation to Length for neighbors 0.08099163536800867
Correlation to Length for income 0.009641611840837548
```

In [595...

```
vectorAssembler = VectorAssembler(inputCols = ['income', 'neighbors', 'looped', 'count'],
vhouse_df = vectorAssembler.transform(house_df)
vhouse_df = vhouse_df.select(['features', 'length'])
vhouse_df.show(3)
```

```
+-----+-----+
|          features|length|
+-----+-----+
|[1.0005,2.0,0.0,1...|    18|
|[1.0,1.0,0.0,1.0,...|    44|
|[2.0,2.0,0.0,1.0,...|     0|
+-----+-----+
only showing top 3 rows
```

Train and Test split

In [596...

```
splits = vhouse_df.randomSplit([0.7, 0.3])
train_df = splits[0]
test_df = splits[1]
```

In [597...

```
test_df.show()
```

```
+-----+-----+
|          features|length|
+-----+-----+
|[0.3,1.0,0.0,1.0,...|    130|
|[0.3,1.0,0.0,1.0,...|    114|
|[0.3,1.0,0.0,1.0,...|    118|
|[0.3,1.0,0.0,1.0,...|     66|
|[0.3,1.0,0.0,1.0,...|     90|
|[0.3,1.0,0.0,1.0,...|     62|
|[0.3,1.0,0.0,1.0,...|     70|
|[0.3,1.0,0.0,1.0,...|    104|
|[0.3,1.0,0.0,1.0,...|     64|
|[0.3,1.0,0.0,1.0,...|     84|
|[0.3,1.0,0.0,1.0,...|    108|
|[0.3,1.0,0.0,1.0,...|     50|
|[0.3,1.0,0.0,1.0,...|     50|
|[0.3,1.0,0.0,1.0,...|     56|
|[0.3,1.0,0.0,1.0,...|     38|
|[0.3,1.0,0.0,1.0,...|     42|
|[0.3,1.0,0.0,1.0,...|     28|
|[0.3,1.0,0.0,1.0,...|     28|
|[0.3,1.0,0.0,1.0,...|     38|
|[0.3,1.0,0.0,1.0,...|     44|
```

```
+-----+-----+
only showing top 20 rows
```

In []:

In [598...

```
from pyspark.ml.regression import LinearRegression
lr = LinearRegression(featuresCol = 'features', labelCol='length', maxIter=10, regPa
lr_model = lr.fit(train_df)
print("Coefficients: " + str(lr_model.coefficients))
print("Intercept: " + str(lr_model.intercept))
```

```
Coefficients: [0.0,0.0,0.0,0.0,0.0,0.9933528371493169]
Intercept: 0.1747814056096046
```

In [599...

```
trainingSummary = lr_model.summary
print("RMSE: %f" % trainingSummary.rootMeanSquaredError)
print("r2: %f" % trainingSummary.r2)
```

```
RMSE: 0.283893
r2: 0.999956
```

In [600...

```
train_df.describe().show()
```

```
+-----+-----+
|summary|          length|
+-----+-----+
|  count|          733966|
|   mean| 26.2941362406433|
| stddev|42.70898842474936|
|   min|              0|
|   max|             144|
+-----+-----+
```

In [601...

```
lr_predictions = lr_model.transform(test_df)
lr_predictions.select("prediction", "length", "features").show(5)
from pyspark.ml.evaluation import RegressionEvaluator
lr_evaluator = RegressionEvaluator(predictionCol="prediction", \
                                   labelCol="length", metricName="r2")
print("R Squared (R2) on test data = %g" % lr_evaluator.evaluate(lr_predictions))
```

```
+-----+-----+-----+
|      prediction|length|          features|
+-----+-----+-----+
| 129.3106502350208|   130|[0.3,1.0,0.0,1.0,...|
|113.41700484063173|   114|[0.3,1.0,0.0,1.0,...|
| 117.390416189229|   118|[0.3,1.0,0.0,1.0,...|
| 65.73606865746453|    66|[0.3,1.0,0.0,1.0,...|
| 89.57653674904813|    90|[0.3,1.0,0.0,1.0,...|
+-----+-----+-----+
only showing top 5 rows
```

```
R Squared (R2) on test data = 0.999956
```

In [602...

```
test_result = lr_model.evaluate(test_df)
print("Root Mean Squared Error (RMSE) on test data = %g" % test_result.rootMeanSquar
```

```
Root Mean Squared Error (RMSE) on test data = 0.284139
```


In [603...

```
print("numIterations: %d" % trainingSummary.totalIterations)
print("objectiveHistory: %s" % str(trainingSummary.objectiveHistory))
trainingSummary.residuals.show()
```

```
numIterations: 11
objectiveHistory: [0.4999999999999999, 0.38122135545559316, 0.04948092617442175, 0.0
29621308549084034, 0.007708081762276442, 0.006428176910368608, 0.006297748096204369,
0.0062975373843313755, 0.006297411204998148, 0.006297335645781342, 0.006297290399106
494]
```

```
+-----+
|          residuals|
```

```
+-----+
```

```
| 0.5431121822641671|
| 0.6760554392778317|
|0.46334622805596837|
| 0.3835802738477696|
| 0.5697008336669|
| 0.6095838107709994|
| 0.6627611135764653|
| 0.6760554392778317|
|0.46334622805596837|
| 0.423463250951869|
| 0.5165235308614342|
| 0.5298178565628007|
| 0.5032292051600677|
| 0.5431121822641671|
| 0.3171086453409373|
| 0.6095838107709994|
| 0.3569916224450367|
| 0.2506370168341121|
|0.33040297104230376|
| 0.5564065079655336|
```

```
+-----+
```

only showing top 20 rows

In [604...

```
predictions = lr_model.transform(test_df)
predictions.select("prediction", "length", "features").show()
```

```
+-----+-----+-----+
|          prediction|length|          features|
```

```
+-----+-----+-----+
```

```
| 129.3106502350208| 130|[0.3,1.0,0.0,1.0,...|
|113.41700484063173| 114|[0.3,1.0,0.0,1.0,...|
| 117.390416189229| 118|[0.3,1.0,0.0,1.0,...|
| 65.73606865746453| 66|[0.3,1.0,0.0,1.0,...|
| 89.57653674904813| 90|[0.3,1.0,0.0,1.0,...|
|61.762657308867254| 62|[0.3,1.0,0.0,1.0,...|
| 69.7094800060618| 70|[0.3,1.0,0.0,1.0,...|
|103.48347646913857| 104|[0.3,1.0,0.0,1.0,...|
| 63.74936298316589| 64|[0.3,1.0,0.0,1.0,...|
| 83.61641972615223| 84|[0.3,1.0,0.0,1.0,...|
|107.45688781773583| 108|[0.3,1.0,0.0,1.0,...|
| 49.84242326307545| 50|[0.3,1.0,0.0,1.0,...|
| 49.84242326307545| 50|[0.3,1.0,0.0,1.0,...|
|55.802540285971354| 56|[0.3,1.0,0.0,1.0,...|
| 37.92218921728365| 38|[0.3,1.0,0.0,1.0,...|
| 41.89560056588092| 42|[0.3,1.0,0.0,1.0,...|
|27.988660845790477| 28|[0.3,1.0,0.0,1.0,...|
|27.988660845790477| 28|[0.3,1.0,0.0,1.0,...|
| 37.92218921728365| 38|[0.3,1.0,0.0,1.0,...|
| 43.88230624017955| 44|[0.3,1.0,0.0,1.0,...|
```

```
+-----+-----+-----+
```

only showing top 20 rows

In [605...

```
test_df
```

Out[605...

```
DataFrame[features: vector, length: int]
```

In [606...

```
from pyspark.ml.regression import DecisionTreeRegressor
dt = DecisionTreeRegressor(featuresCol='features', labelCol='length')
dt_model = dt.fit(train_df)
dt_predictions = dt_model.transform(test_df)
dt_evaluator = RegressionEvaluator(
    labelCol="length", predictionCol="prediction", metricName="rmse")
rmse = dt_evaluator.evaluate(dt_predictions)
print("Root Mean Squared Error (RMSE) on test data = %g" % rmse)
```

```
Root Mean Squared Error (RMSE) on test data = 2.39521
```

In []:

In []: